


SIGNLL Meeting 3



Join Discord!





Agenda (9/26)

- Join the Discord (hopefully it works this time)
- Naive Bayes!
- Start thinking of project ideas



Naive Bayes!

- A probabilistic classifier!
- Out of all classes $c \in C$, it returns the class \hat{c} which is most likely.
- Uses:
 - Sentiment analysis!
 - Spam detection!
 - More generally, **text categorization**:
 - Tasks in which we label or categorize entire documents
- Let's look at Naive Bayes in the context of sentiment analysis for the next few slides: our classes will be either positive or negative.



Bag of Words

Document	the	cat	sat	in	hat	with
<i>the cat sat</i>	1	1	1	0	0	0
<i>the cat sat in the hat</i>	2	1	1	1	1	0
<i>the cat with the hat</i>	2	1	0	0	1	1

What makes Naive Bayes “naive” is that we don’t care about word order. We can represent a document as an unordered **bag of words**, as shown above.



What is Naive Bayes?

Bayes' rule:

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$



A bunch of stats and algebra later...

$P(c)$ is the prior probability of a class among all documents.

$P(w_i | c)$ is the probability of a particular word given that a document is of a certain class (formula for this on the next slide).

$$c_{NB} = \arg \max_{c \in C} \log P(c) + \sum_{i \in positions} \log P(w_i | c)$$



Laplace Smoothing

The added +1 and +|V| in the numerator and denominator are a result of Laplace smoothing. What if the word “great” appears in a negative document but not a positive one, in our data? Then $P(\text{great}|\text{positive}) = 0$, so the entire probability would be 0. Not ideal.

$$P(w_i|c) = \frac{\text{count}(w_i, c) + 1}{\sum_w \text{count}(w, c) + |V|}$$



Try it out!

- Workshop: Use Naive Bayes and NLTK to classify tweets as either coming from Joe Biden or Kanye West



ye ✓

@kanyewest



Joe Biden ✓

@JoeBiden